


Attention Refinement




Base View

Patch Embedding and Norm

Multi-head Attention Layer



$$\hat{A}^{(b)}(x) = I + E_h(A^{(b)}(x) \text{diag}(W^{(b)}(x)))^+$$



...

Last-2 Average Rollout


Final View



Attention-Guided Multi-View Augmentation



Image Encoder




Prompt Tuning

[class 1 ... class C]




Text Encoder


VT-Based Ensemble



$TV(A(x_i)) \Rightarrow \omega_i$



Final Output



Entropy Minimization